ASSIGNMENT 1:
CIFAR100 IMAGE CLASSIFICATION WITH KERAS

INTRODUCTION

CIFAR-100

- *100 CLASSES WITH 600 IMAGES EACH*

Process:

- Load and preprocess dataset
  - normalizing pixel values and converting class labels into one-hot encoded vectors
- construct a CNN architecture
  - Used established model e.g. VGG-16
  - trained it on the preprocessed dataset

Training:

- Determined loss and accuracy for training and validation

Experiment:

- Trained different models using base vgg-16 model, extra block of layers, use SGD optimization, and a smaller architecture

MODEL ARCHITECTURE

- VGG-16 :
- Input Layer:
- - Input dimensions: (32, 32, 3) for CIFAR-100 dataset.
- Convolutional Layers:
- - 2 convolutional layers with 64 filters each, kernel size of (3, 3), and ReLU activation function.
- - 2 convolutional layers with 128 filters each, kernel size of (3, 3), and ReLU activation function.
- - 3 convolutional layers with 256 filters each, kernel size of (3, 3), and ReLU activation function.
- - 3 convolutional layers with 512 filters each, kernel size of (3, 3), and ReLU activation function.
- Max Pooling Layers:
- - After every two convolutional layers, there's a max-pooling layer with a pool size of (2, 2) and stride of (2, 2).
- Fully Connected Layers:
- - After the convolutional layers, there are three fully connected layers with ReLU activation function.
- - 2 fully connected layers have 4096 units each.
- - The last fully connected layer has 100 units with softmax activation, corresponding to the 100 classes in the CIFAR-100 dataset.
- Regularization:
- - Dropout layers with a dropout rate of 0.5 are added after the first two fully connected layers.
- Batch Normalization:
- - Batch normalization layers are added after every max-pooling layer to stabilize and accelerate the training process.

MODELS

Model 1 :

- Base model
- Adam optimizer
- Epoch = 20
- Batch = 128

Model 2 :

- Another block of 3 convolutional layers with 512 filters each (from base model)
- Adam optimizer
- Epoch = 20
- Batch = 128

Model 3 :

- same architecture as model 1
- Use SGD optimizer
- Epoch = 20
- Batch = 128

Model 3 :

- Smaller network (Input Shape: (32, 32, 3)
- Conv2D(32, (3, 3)) -> MaxPooling2D -> Dropout
- Conv2D(64, (3, 3)) -> MaxPooling2D -> Dropout
- Flatten
- Dense(512) -> Dropout
- Dense(num_classes, softmax activation)
- Adam optimizer, Epoch = 38, Batch = 128

KERAS VISUAL

Model 1: Base model

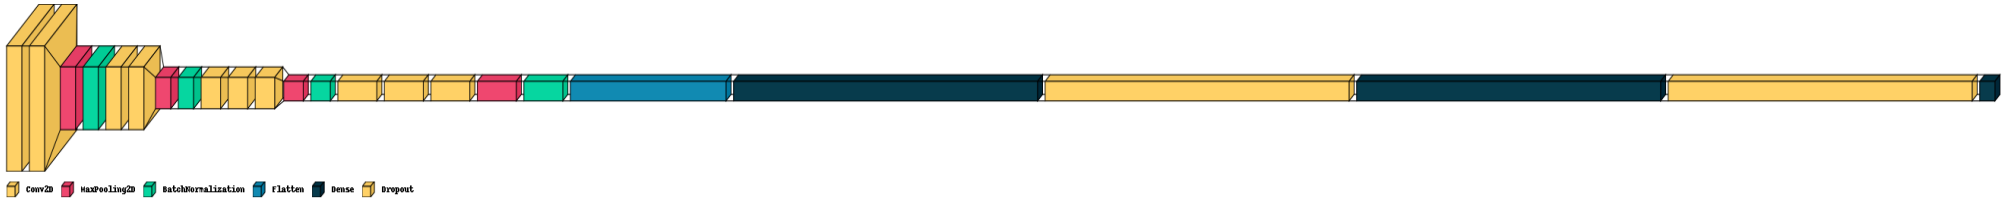

Model 2: With another block of 3 convolutional layers (512 filters)

Model 3: same as model 2

Model 4: With another block of 3 convolutional layers(512 filters)

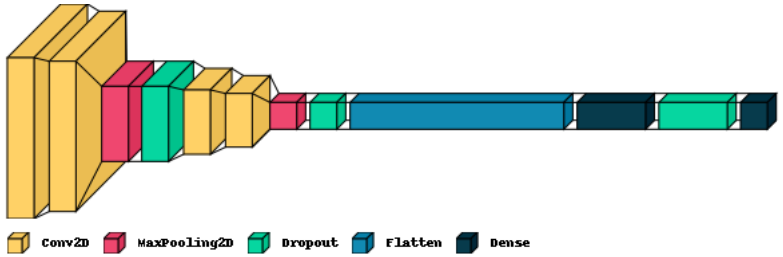

MODEL 1

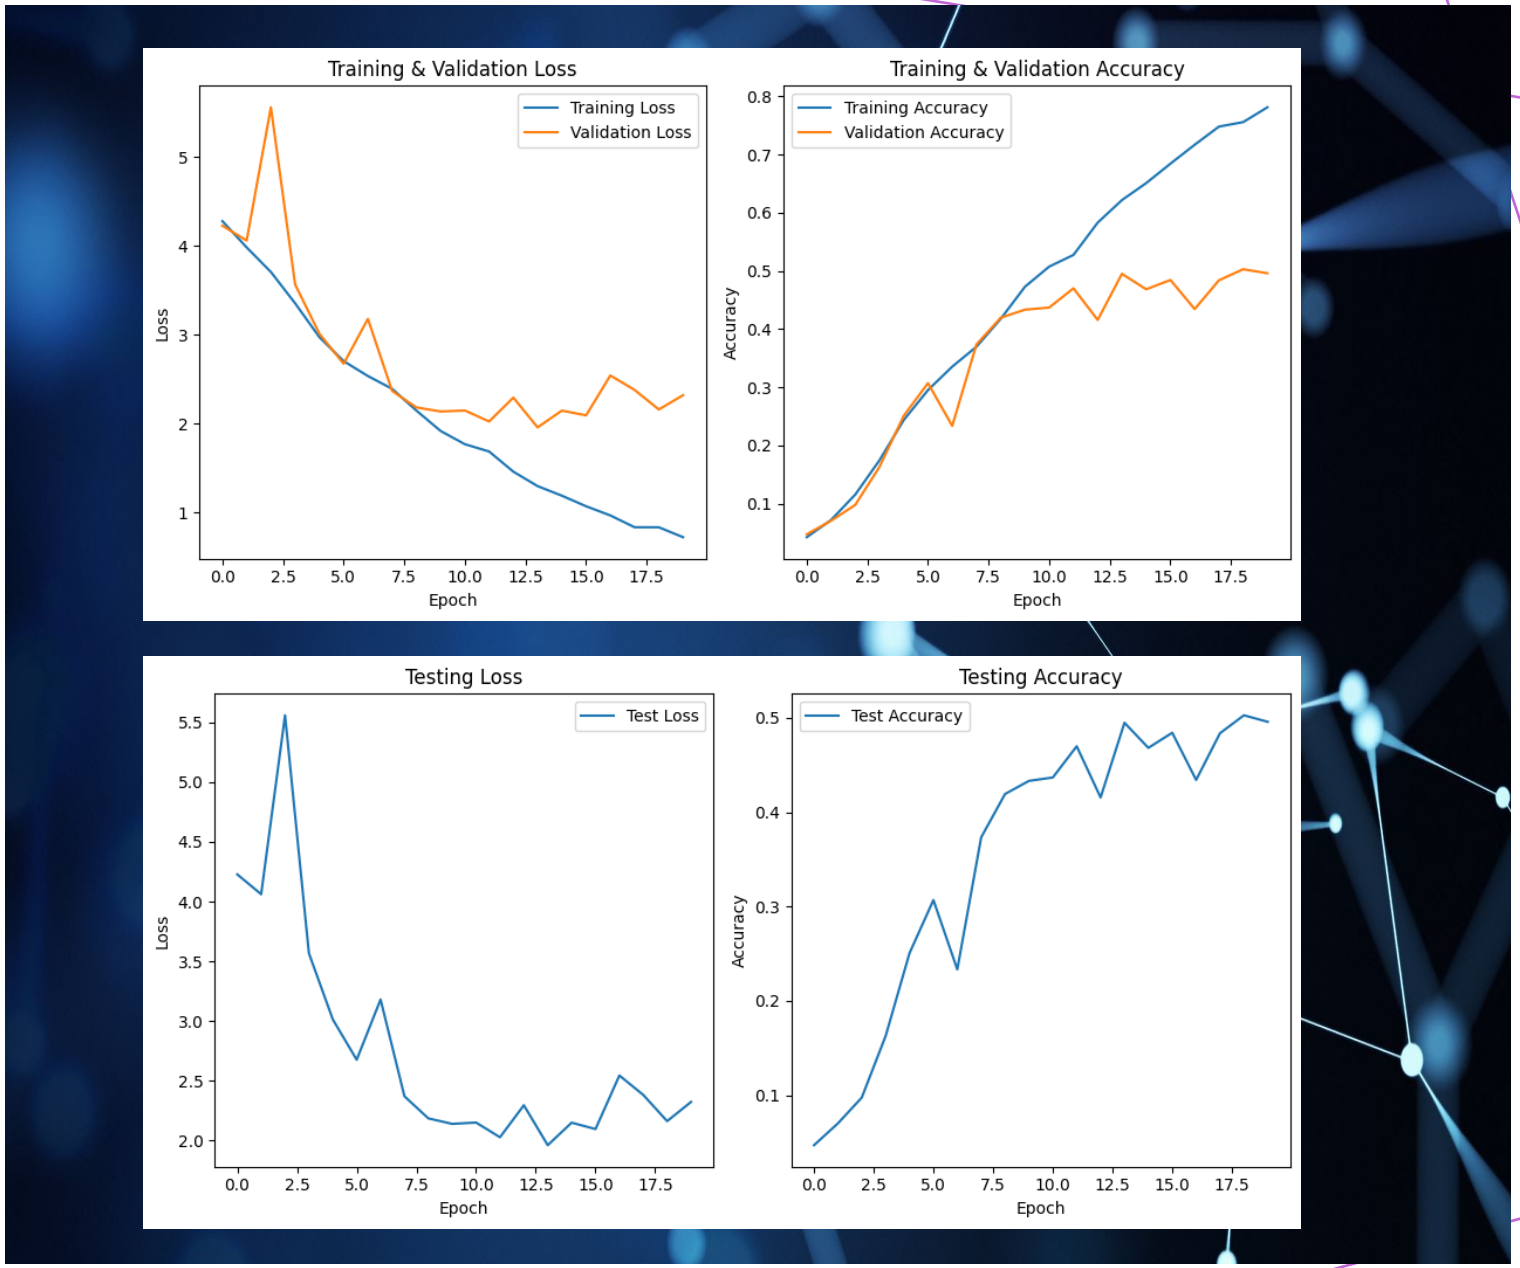
Training & Validation Loss
Training Loss
Validation Loss
Loss
Epoch
Training & Validation Accuracy
Training Accuracy
Validation Accuracy
Accuracy
Epoch
Testing Loss
Test Loss
Loss
Epoch
Testing Accuracy
Test Accuracy
Accuracy
Epoch

MODEL 2

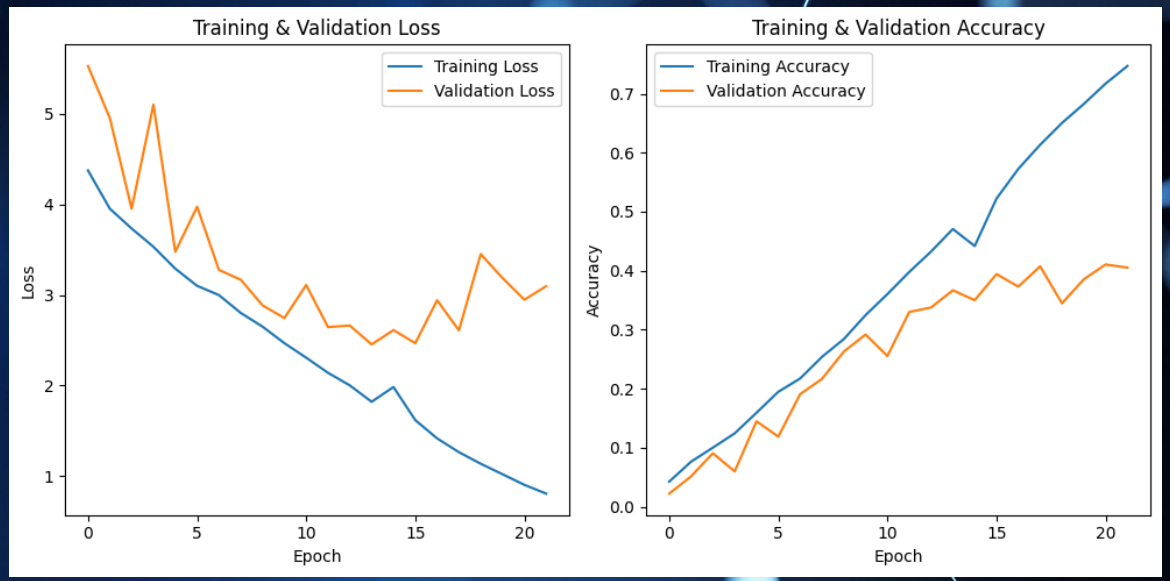
Training & Validation Loss
Training Loss
Validation Loss
Loss
Epoch
Training & Validation Accuracy
Training Accuracy
Validation Accuracy
Accuracy
Epoch

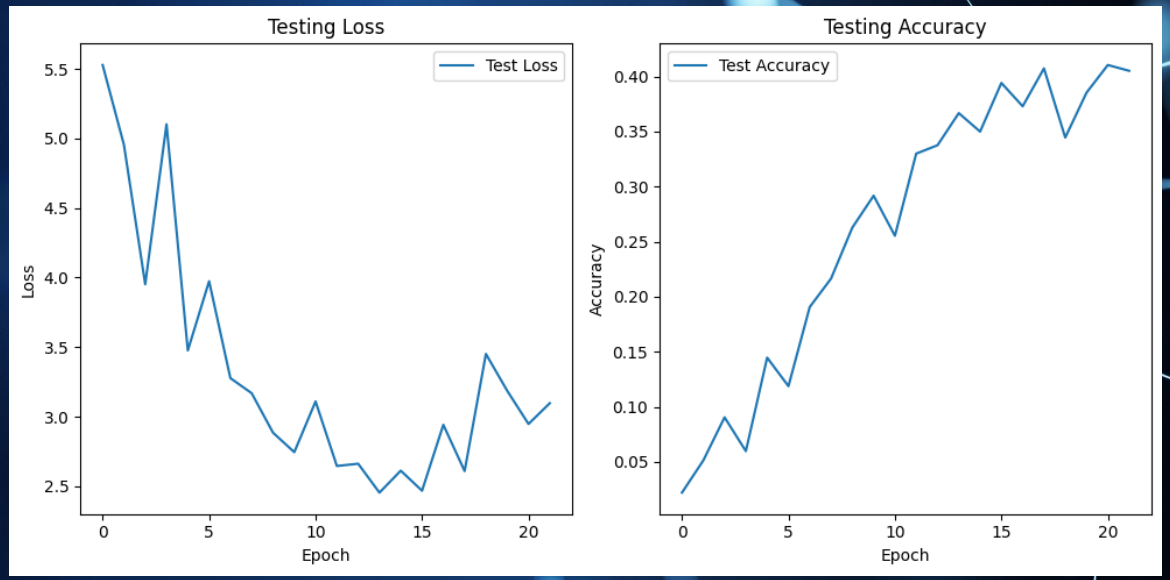
Testing Loss
Test Loss
Loss
Epoch
Testing Accuracy
Test Accuracy
Accuracy
Epoch

MODEL 3

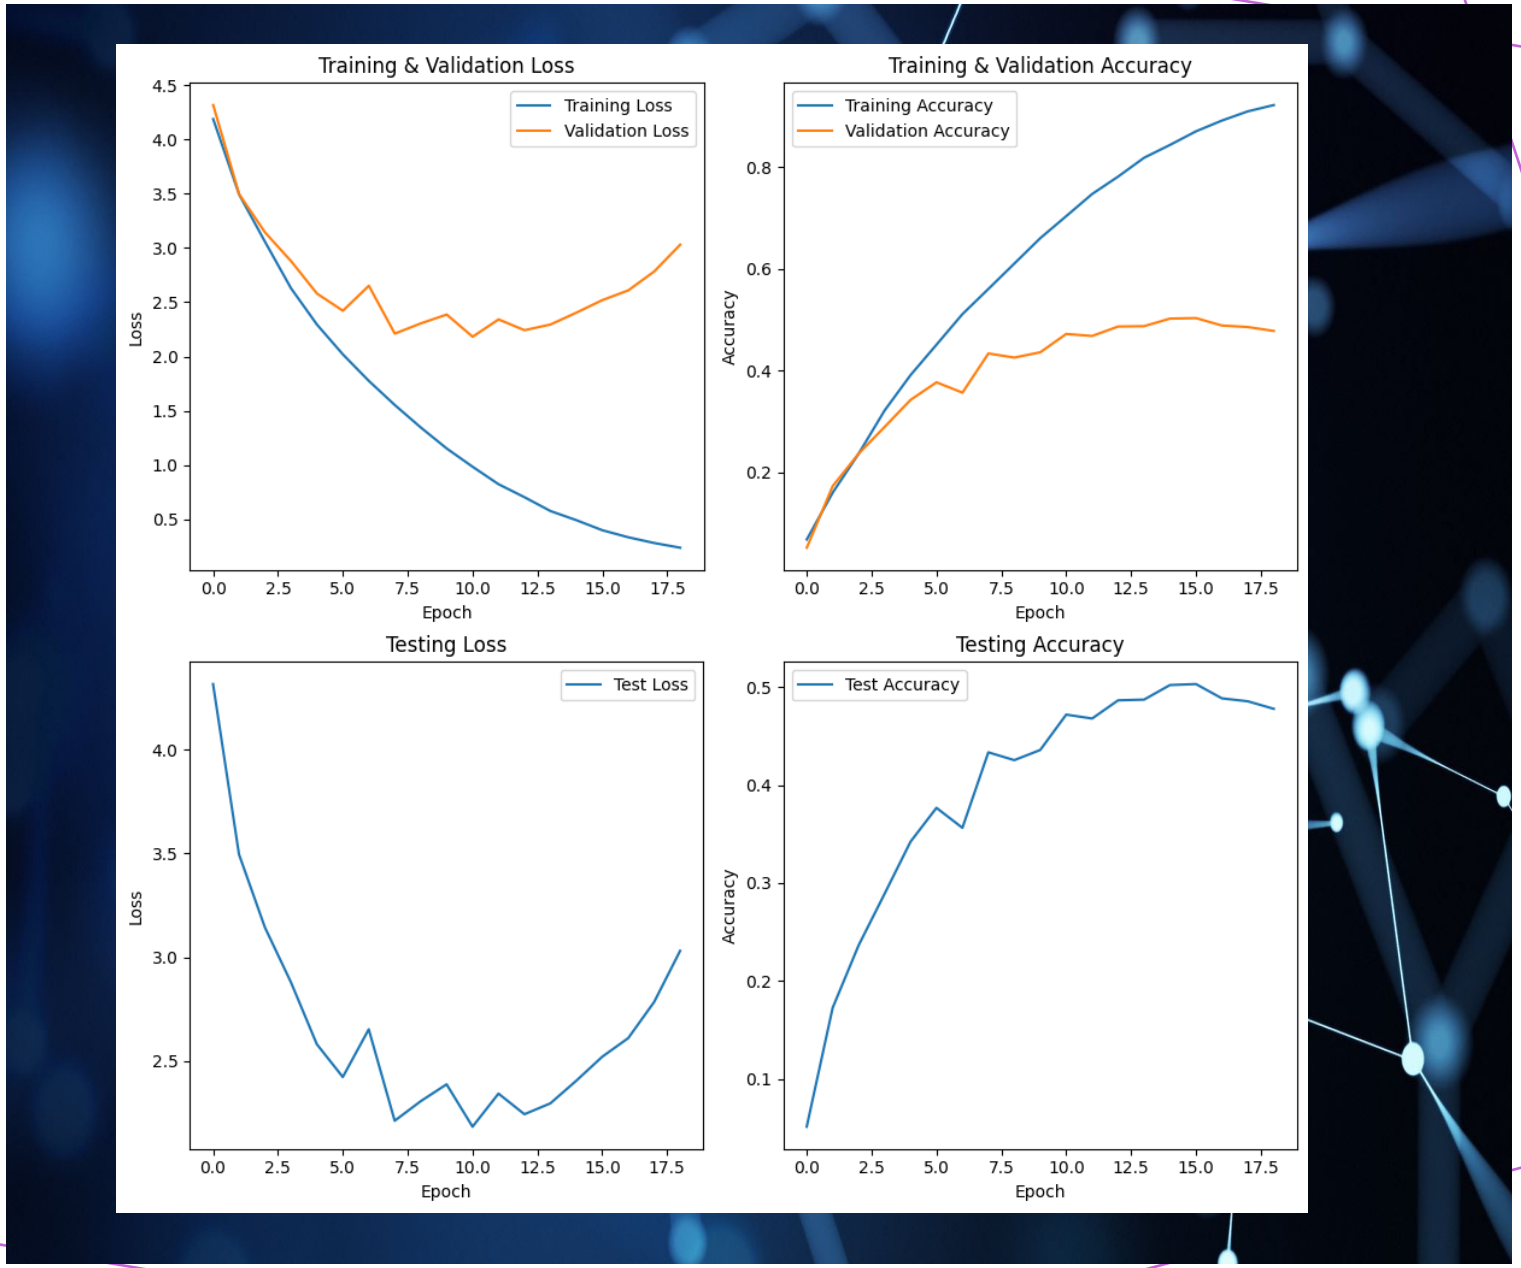
Training & Validation Loss
Training Loss
Validation Loss
Loss
Epoch
Training & Validation Accuracy
Training Accuracy
Validation Accuracy
Accuracy
Epoch
Testing Loss
Test Loss
Loss
Epoch
Testing Accuracy
Test Accuracy
Accuracy
Epoch

MODEL 4

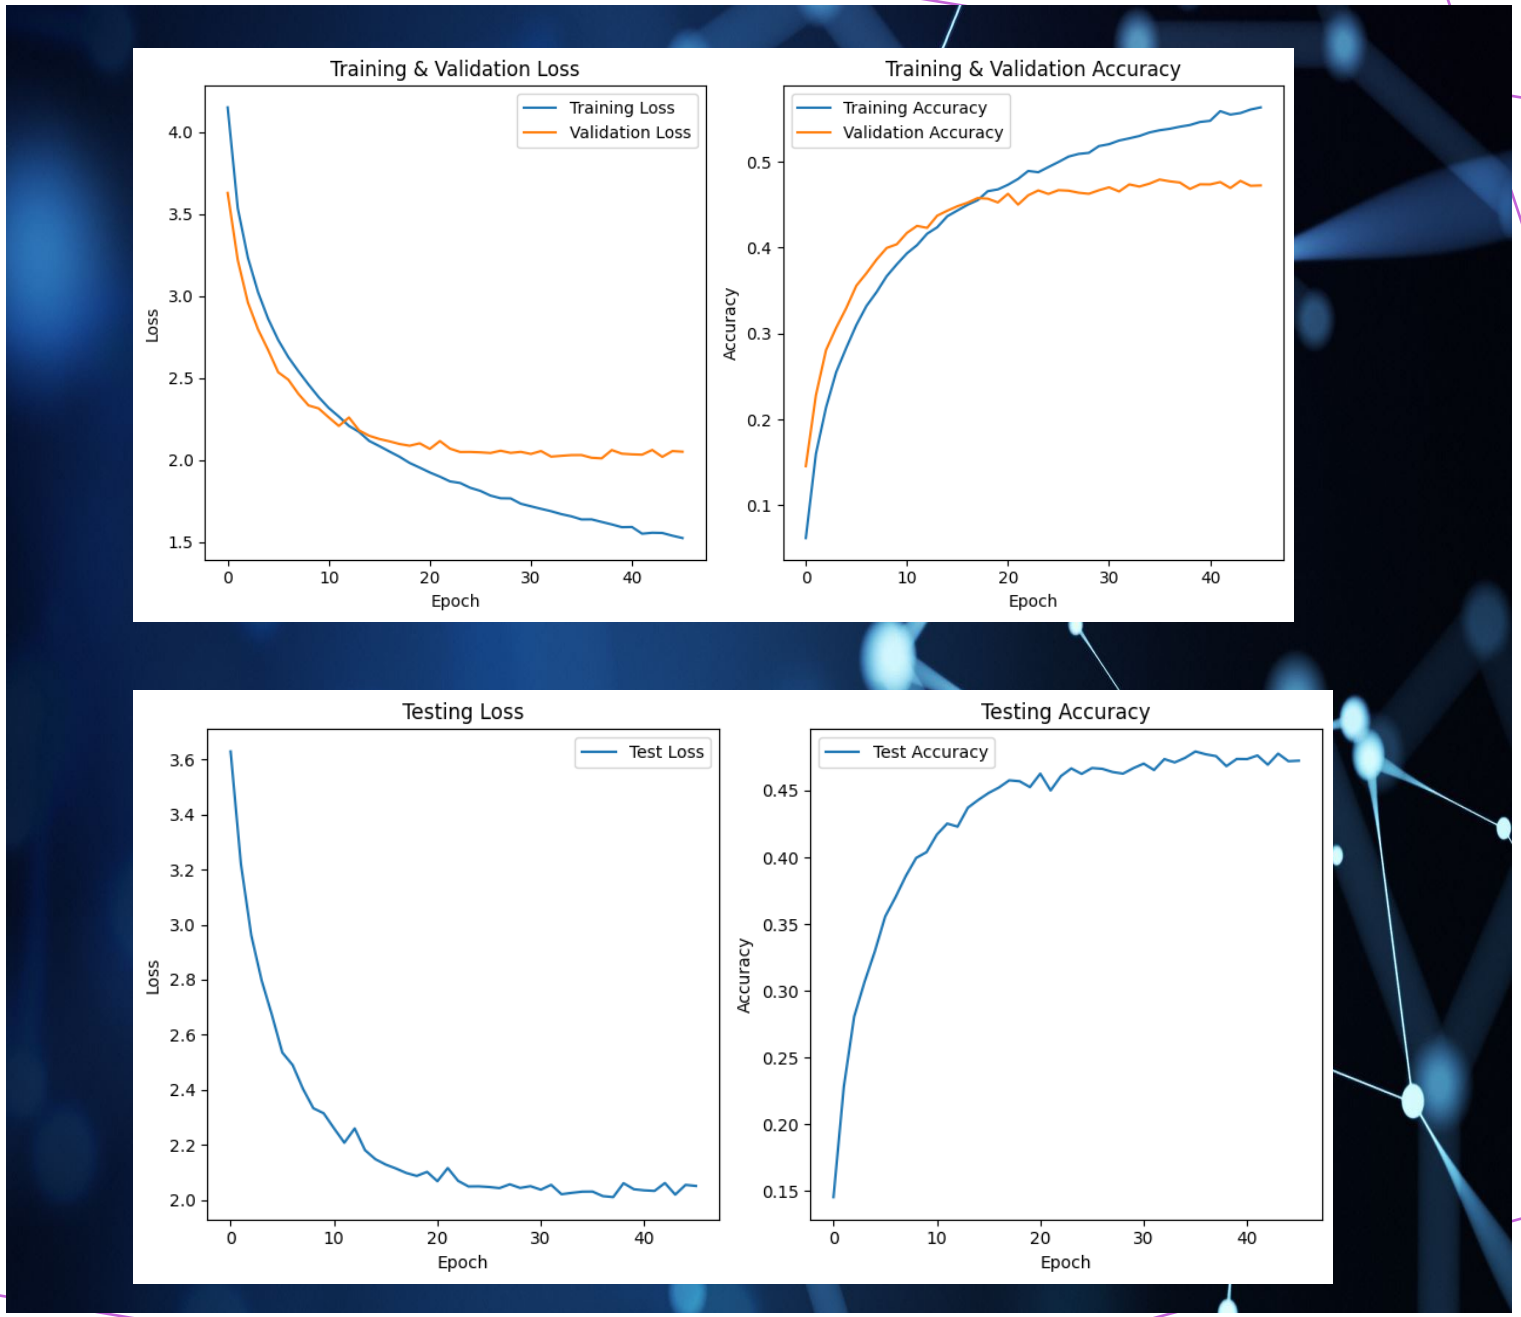
Training & Validation Loss
Training Loss
Validation Loss
Loss
Epoch
Training & Validation Accuracy
Training Accuracy
Validation Accuracy
Accuracy
Epoch
Testing Loss
Test Loss
Loss
Epoch
Testing Accuracy
Test Accuracy
Accuracy
Epoch

MODEL RESULTS

- Model 1 (BEST)
  - Training accuracy: 0.78
  - Test loss: 2.32
  - Test accuracy: 0.50

- Model 2
  - Training accuracy: 0.47
  - Test loss: 2.45
  - Test accuracy: 0.37

- Model 3
  - Training accuracy: 0.70
  - Test loss: 2.18
  - Test accuracy: 0.47

- Model 4
  - Training accuracy: 0.54
  - Test loss: 2.01
  - Test accuracy: 0.48

IMAGE LABELING

- display of test images, only 2 correctly labeled

Predicted: lizard
True: mountain

Predicted: squirrel
True: forest

Predicted: otter
True: seal

Predicted: mushroom
True: mushroom

Predicted: sea
True: sea

Predicted: bee
True: tulip

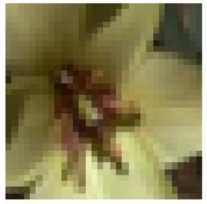

Predicted: kangaroo
True: camel

Predicted: chimpanzee
True: butterfly

Predicted: sea
True: cloud

IMAGE LABELING

- Selection of correctly labeled images

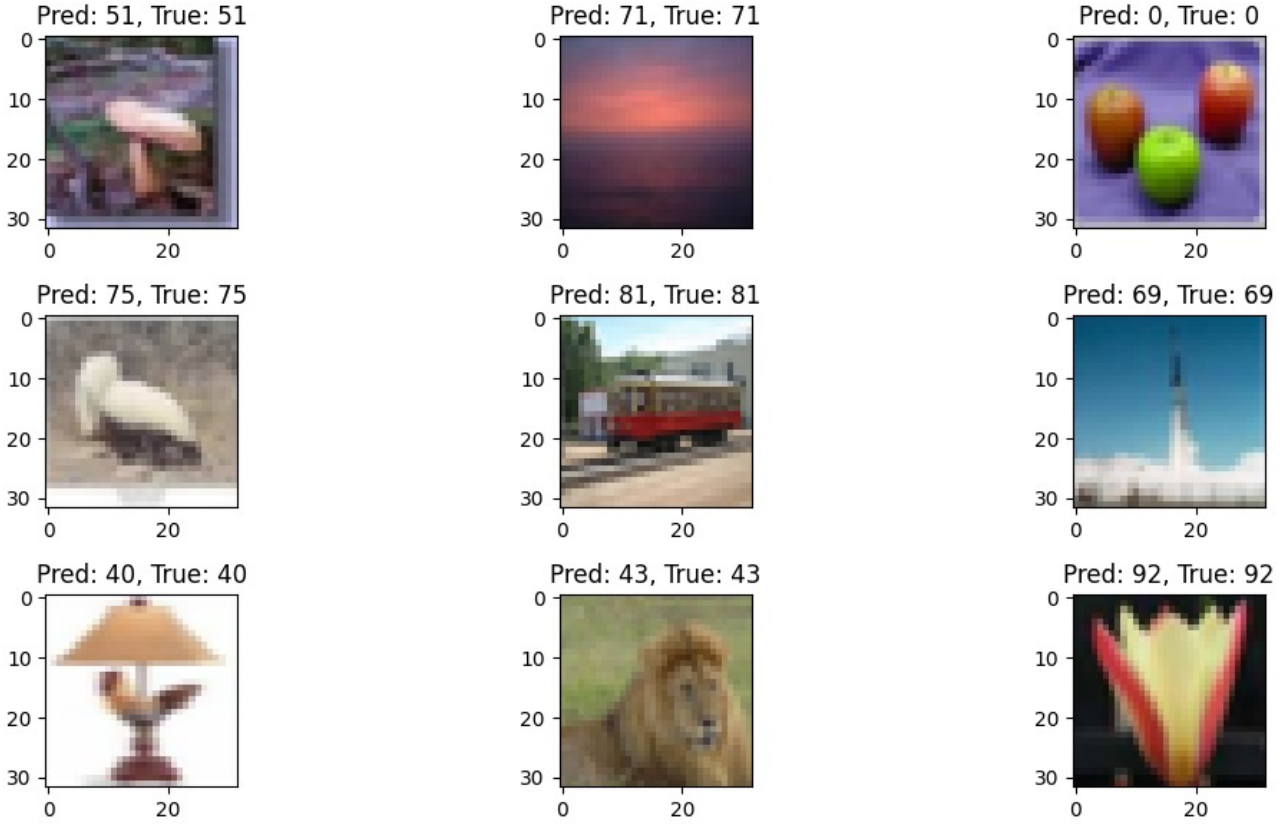

IMAGE LABELING

- Selection of incorrectly labeled images

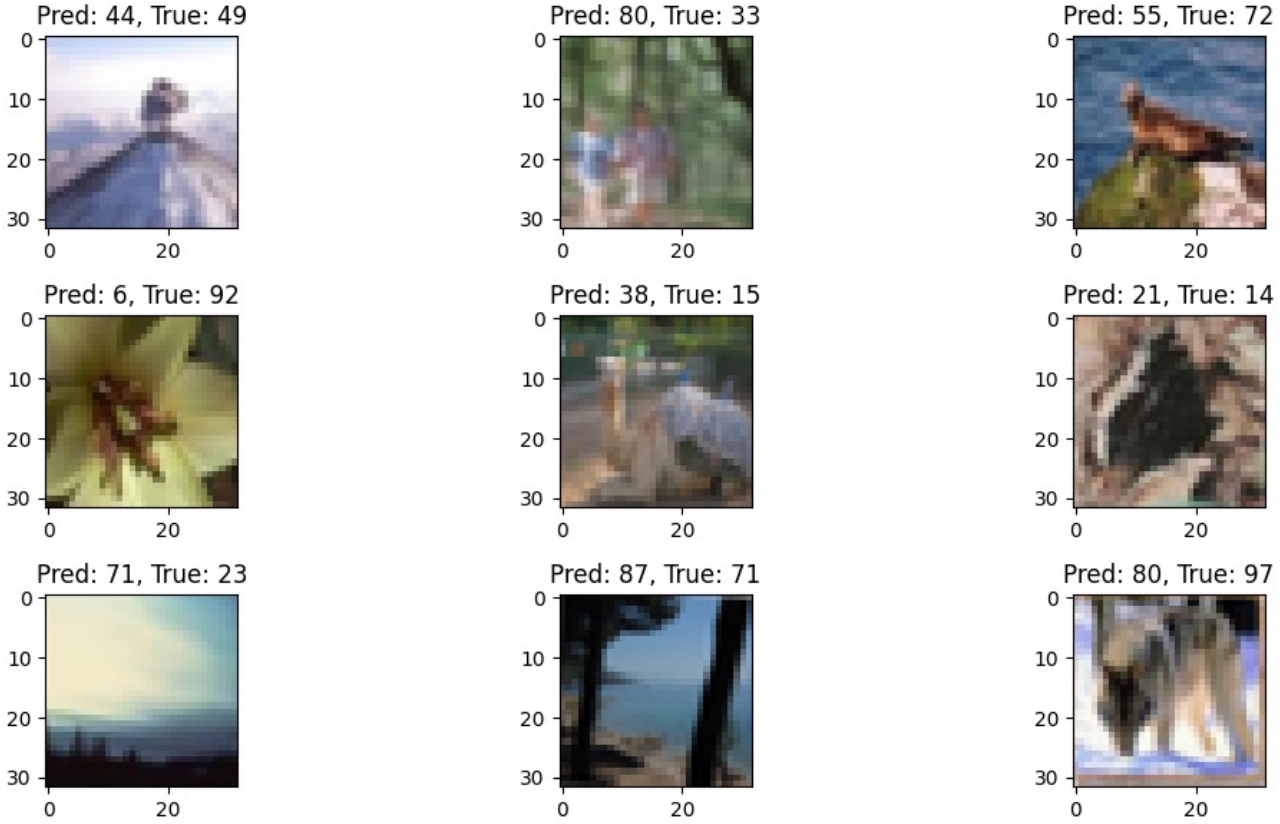

CONCLUSION

- Model 1 had the best test accuracy compared to the other models with a test accuracy of 0.50.
- I could improve the training and testing accuracy using another optimizer or architecture.