

ASSIGNMENT 1:
CIFAR100 IMAGE
CLASSIFICATION WITH
KERAS



INTRODUCTION

CIFAR-100

- 100 CLASSES WITH 600 IMAGES EACH

Process:

- Load and preprocess dataset
 - normalizing pixel values and converting class labels into one-hot encoded vectors
- construct a CNN architecture
 - Used established model e.g. VGG-16
 - trained it on the preprocessed dataset

Training:

- Determined loss and accuracy for training and validation

Experiment:

- Trained different models with batch normalization, no batch normalization, and added another block (4 instead of 3)

MODEL ARCHITECTURE

VGG-16 :

Input Layer:

- Input dimensions: (32, 32, 3) for CIFAR-100 dataset.

Convolutional Layers:

- 2 convolutional layers with 64 filters each, kernel size of (3, 3), and ReLU activation function.
- 2 convolutional layers with 128 filters each, kernel size of (3, 3), and ReLU activation function.
- 3 convolutional layers with 256 filters each, kernel size of (3, 3), and ReLU activation function.
- 3 convolutional layers with 512 filters each, kernel size of (3, 3), and ReLU activation function.

Max Pooling Layers:

- After every two convolutional layers, there's a max-pooling layer with a pool size of (2, 2) and stride of (2, 2).

Fully Connected Layers:

- After the convolutional layers, there are three fully connected layers with ReLU activation function.
- 2 fully connected layers have 4096 units each.
- The last fully connected layer has 100 units with softmax activation, corresponding to the 100 classes in the CIFAR-100 dataset.

Regularization:

- Dropout layers with a dropout rate of 0.5 are added after the first two fully connected layers.

Batch Normalization:

- Batch normalization layers are added after every max-pooling layer to stabilize and accelerate the training process.

MODEL

Model 1 :

- No batch normalization
- Epoch = 20
- Batch = 164

Model 2 :


- Batch normalization
- Epoch = 20
- Batch = 164

Model 3 :


- Another block of 3 convolutional layers with 512 filters each
- Epoch = 20
- Batch = 164

KERAS VISUAL

Without batch normalization





With batch normalization





With another block of 3 convolutional layers(512 filters)





MODEL 1



MODEL 2



MODEL 3



RESULTS

- Best
- Model 2
 - training accuracy: 0.62
 - Test loss: 1.98
 - Test accuracy: 0.49
- Model 1
 - training accuracy: 0.0084
 - Test loss: 4.61
 - Test accuracy: 0.01
- Model 3
 - training accuracy: 0.54
 - Test loss: 2.94
 - Test accuracy: 0.37

IMAGE LABELING

- display of
test images,
only 2
correctly
labeled

Predicted: otter
True: mountain



Predicted: mushroom
True: mushroom




Predicted: otter
True: camel




Predicted: kangaroo
True: forest



Predicted: sea
True: sea




Predicted: beetle
True: butterfly



Predicted: beaver
True: seal



Predicted: bee
True: tulip



Predicted: sea
True: cloud



IMAGE LABELING

- Selection of
correctly labeled
images












IMAGE LABELING

- Selection of
incorrectly labeled
images





















IMAGE LABELING

- Selection of
incorrectly labeled
images



CONCLUSION

- Model 2 had the best test accuracy compared to the other models with an accuracy of 0.49.
- Model 2 with epoch 20 and batch size 164 was stopped early to prevent overfitting which ended up with a training accuracy of 0.62. If the early stop was not there, the training for epoch 20 was up to 0.82. I could improve the training and testing accuracy using another optimizer or architecture.